

Handout Week 3: Pandas for Tabular Data

Dr Ian (Yinglong) He

[Click to download the Jupyter Notebook files \(.ipynb\)](#)

based on the notes by Jake VanderPlas

Contents

1	Introduction	2
2	Constructing DataFrame objects	3
2.1	From a two-dimensional NumPy array	3
2.2	From a dictionary	3
3	Writing DataFrame to a CSV File	4
4	Data Selection in DataFrame	4
4.1	DataFrame as a dictionary	4
4.2	DataFrame as two-dimensional array	5
4.2.1	Indexers: loc and iloc	5
4.3	Additional indexing conventions	7
5	Operating on Null Values	7
5.1	Detecting null values	8
5.2	Dropping null values	8
5.3	Filling null values	8
6	Combining Datasets: Concat	9
6.1	Ignoring the index	10
6.2	Concatenation with joins	11
7	Combining Datasets: Merge	12
7.1	Categories of Joins	12
7.1.1	One-to-one joins	12
7.1.2	Many-to-one joins	13
7.1.3	Many-to-many joins	14
7.2	Specification of the Merge Key	15
7.2.1	The on keyword	15
7.2.2	The left_on and right_on keywords	15
8	Aggregation and Grouping	16
8.1	Simple Aggregation in Pandas	17
8.2	GroupBy: Split, Apply, Combine	18
8.2.1	Split, apply, combine	18
9	Reading and Writing Files	20

1 Introduction

As we saw, NumPy's `ndarray` data structure provides essential features for the type of **clean, well-organized data** typically seen in numerical computing tasks. While it serves this purpose very well, its limitations become clear when we need more flexibility (e.g., attaching labels to data, working with missing data, etc.) and when attempting operations that do not map well to element-wise broadcasting (e.g., groupings, pivots, etc.), each of which is an important piece of analyzing the less structured data available in many forms in the world around us.

Pandas, and in particular its `Series` and `DataFrame` objects, builds on the NumPy array structure. If a `Series` is an analog of a one-dimensional array with flexible indices, a `DataFrame` is an analog of a two-dimensional array with both flexible row indices and flexible column names. `DataFrames` are essentially multidimensional arrays with attached row and column labels, and often with **heterogeneous types and/or missing data**. As well as offering a convenient storage interface for labeled data, Pandas implements a number of powerful data operations familiar to users of both database frameworks and spreadsheet programs.

In this handout, we will focus on the mechanics of using `DataFrame`, and related structures effectively. We will use examples drawn from real datasets where appropriate, but these examples are not necessarily the focus. Details on this installation can be found in the [Pandas documentation](#). If you used the Anaconda stack, you already have Pandas installed.

Once Pandas is installed, you can import it and check the version:

```
[1]: import pandas
      pandas.__version__
```

```
[1]: '1.5.1'
```

Just as we generally import NumPy under the alias `np`, we will import Pandas under the alias `pd`:

```
[2]: import pandas as pd
```

This import convention will be used throughout the remainder of this book.

For convenience, we'll use the same `display` magic function that we've seen in previous sections:

```
[3]: import numpy as np
      import pandas as pd

      class display(object):
          """Display HTML representation of multiple objects"""
          template = """<div style="float: left; padding: 10px;">
<p style="font-family:'Courier New', Courier, monospace'>{0}</p>{1}
</div>"""
          def __init__(self, *args):
              self.args = args

          def _repr_html_(self):
              return '\n'.join(self.template.format(a, eval(a)._repr_html_())
                              for a in self.args)

          def __repr__(self):
```

```
return '\n\n'.join(a + '\n' + repr(eval(a))
                    for a in self.args)
```

2 Constructing DataFrame objects

A Pandas DataFrame can be constructed in a variety of ways. Here we'll give several examples.

2.1 From a two-dimensional NumPy array

Given a two-dimensional array of data, we can create a DataFrame with any specified column and index names. If omitted, an integer index will be used for each:

```
[4]: pd.DataFrame(np.random.rand(3, 2),
                  columns=['foo', 'bar'],
                  index=['a', 'b', 'c'])
```

```
[4]:
```

	foo	bar
a	0.245258	0.496904
b	0.316460	0.871386
c	0.918690	0.193112

2.2 From a dictionary

Any dictionary can be made into a DataFrame:

```
[5]: data_dict = {
      'state': ['California', 'Texas', 'New York', 'Florida', 'Illinois'],
      'area': [423967, 695662, 141297, 170312, 149995],
      'population': [38332521, 26448193, 19651127, 19552860, 12882135]}

df = pd.DataFrame(data_dict)
df
```

```
[5]:
```

	state	area	population
0	California	423967	38332521
1	Texas	695662	26448193
2	New York	141297	19651127
3	Florida	170312	19552860
4	Illinois	149995	12882135

We can set the index to become one of the columns:

```
[6]: df = df.set_index('state')
df
```

```
[6]:
```

	area	population
state		
California	423967	38332521
Texas	695662	26448193
New York	141297	19651127
Florida	170312	19552860
Illinois	149995	12882135

The `DataFrame` have an `index` attribute that gives access to the index labels:

```
[7]: df.index
```

```
[7]: Index(['California', 'Texas', 'New York', 'Florida', 'Illinois'],
      dtype='object', name='state')
```

Additionally, the `DataFrame` and `Series` have a `columns` attribute, which is an `Index` object holding the column labels:

```
[8]: df.columns
```

```
[8]: Index(['area', 'population'], dtype='object')
```

3 Writing DataFrame to a CSV File

We can store a `DataFrame` object into a csv file using the `to_csv` method:

```
[9]: df.to_csv('my_dataframe.csv')
```

4 Data Selection in DataFrame

In the previous handout, we looked in detail at methods and tools to access, set, and modify values in NumPy arrays. These included indexing (e.g., `arr[2, 1]`), slicing (e.g., `arr[:, 1:5]`), masking (e.g., `arr[arr > 0]`), fancy indexing (e.g., `arr[0, [1, 5]]`), and combinations thereof (e.g., `arr[:, [1, 5]]`). Here we'll look at similar means of accessing and modifying values in Pandas `DataFrame` objects. If you have used the NumPy patterns, the corresponding patterns in Pandas will feel very familiar, though there are a few quirks to be aware of.

4.1 DataFrame as a dictionary

The first analogy we will consider is the `DataFrame` as a dictionary of related `Series` objects. Let's return to our example of areas and populations of states:

```
[10]: df['area']
```

```
[10]: state
      California    423967
      Texas        695662
      New York     141297
      Florida      170312
      Illinois     149995
      Name: area, dtype: int64
```

Equivalently, we can use attribute-style access with column names that are strings:

```
[11]: df.area
```

```
[11]: state
      California    423967
      Texas        695662
      New York     141297
      Florida      170312
```

```
Illinois      149995
Name: area, dtype: int64
```

This dictionary-style syntax can also be used to modify the object, in this case adding a new column:

```
[12]: df['density'] = df['population'] / df['area']
df
```

```
[12]:
```

	area	population	density
state			
California	423967	38332521	90.413926
Texas	695662	26448193	38.018740
New York	141297	19651127	139.076746
Florida	170312	19552860	114.806121
Illinois	149995	12882135	85.883763

This shows a preview of the straightforward syntax of element-by-element arithmetic between `Series` objects.

4.2 DataFrame as two-dimensional array

As mentioned previously, we can also view the `DataFrame` as an enhanced two-dimensional array. We can examine the raw underlying data array using the `values` attribute:

```
[13]: df.values
```

```
[13]: array([[4.23967000e+05, 3.83325210e+07, 9.04139261e+01],
             [6.95662000e+05, 2.64481930e+07, 3.80187404e+01],
             [1.41297000e+05, 1.96511270e+07, 1.39076746e+02],
             [1.70312000e+05, 1.95528600e+07, 1.14806121e+02],
             [1.49995000e+05, 1.28821350e+07, 8.58837628e+01]])
```

With this picture in mind, many familiar array-like observations can be done on the `DataFrame` itself. For example, we can transpose the full `DataFrame` to swap rows and columns:

```
[14]: df.T
```

```
[14]: state      California      Texas      New York      Florida \
area      4.239670e+05  6.956620e+05  1.412970e+05  1.703120e+05
population 3.833252e+07  2.644819e+07  1.965113e+07  1.955286e+07
density    9.041393e+01  3.801874e+01  1.390767e+02  1.148061e+02

state      Illinois
area      1.499950e+05
population 1.288214e+07
density    8.588376e+01
```

4.2.1 Indexers: `loc` and `iloc`

For array-style indexing, Pandas uses the `loc` and `iloc` indexers:

- `loc` attribute allows indexing and slicing that always references the explicit (label based) index.

- `iloc` attribute allows indexing and slicing that always references the implicit (integer-location based) index.

Using the `iloc` indexer, we can index the underlying array as if it is a simple NumPy array, but the `DataFrame` index and column labels are maintained in the result:

```
[15]: df.iloc[:3, :2]
```

```
[15]:
```

	area	population
state		
California	423967	38332521
Texas	695662	26448193
New York	141297	19651127

Similarly, using the `loc` indexer we can index the underlying data using the explicit index and column names:

```
[16]: df.loc['Illinois', : 'population']
```

```
[16]:
```

	area	population
state		
California	423967	38332521
Texas	695662	26448193
New York	141297	19651127
Florida	170312	19552860
Illinois	149995	12882135

Any of the familiar NumPy-style data access patterns can be used within these indexers. For example, in the `loc` indexer we can combine masking and fancy indexing as in the following:

```
[17]: df.loc[df['density'] > 100, ['population', 'density']]
```

```
[17]:
```

	population	density
state		
New York	19651127	139.076746
Florida	19552860	114.806121

Any of these indexing conventions may also be used to set or modify values; this is done in the standard way that you might be accustomed to from working with NumPy:

```
[18]: df.iloc[0, 2] = 90
df
```

```
[18]:
```

	area	population	density
state			
California	423967	38332521	90.000000
Texas	695662	26448193	38.018740
New York	141297	19651127	139.076746
Florida	170312	19552860	114.806121
Illinois	149995	12882135	85.883763

To build up your fluency in Pandas data manipulation, I suggest spending some time with a simple `DataFrame` and exploring the types of indexing, slicing, masking, and fancy indexing that are allowed by these various indexing approaches.

4.3 Additional indexing conventions

There are a couple extra indexing conventions that might seem at odds with the preceding discussion, but nevertheless can be very useful in practice. First, while *indexing* refers to columns, *slicing* refers to rows:

```
[19]: df['Florida':'Illinois']
```

```
[19]:
```

	area	population	density
state			
Florida	170312	19552860	114.806121
Illinois	149995	12882135	85.883763

Such slices can also refer to rows by number rather than by index:

```
[20]: df[1:3]
```

```
[20]:
```

	area	population	density
state			
Texas	695662	26448193	38.018740
New York	141297	19651127	139.076746

Similarly, direct masking operations are also interpreted row-wise rather than column-wise:

```
[21]: df[df['density'] > 100]
```

```
[21]:
```

	area	population	density
state			
New York	141297	19651127	139.076746
Florida	170312	19552860	114.806121

These two conventions are syntactically similar to those on a NumPy array, and while these may not precisely fit the mold of the Pandas conventions, they are nevertheless quite useful in practice.

5 Operating on Null Values

The difference between data found in many tutorials and data in the real world is that real-world data is rarely clean and homogeneous. In particular, many interesting datasets will have some amount of data missing. To make matters even more complicated, different data sources may indicate missing data in different ways.

Pandas treats `None` and `NaN` as essentially interchangeable for indicating missing or null values. To facilitate this convention, there are several useful methods for detecting, removing, and replacing null values in Pandas data structures. They are:

- `isnull()`: Generate a boolean mask indicating missing values
- `notnull()`: Opposite of `isnull()`
- `dropna()`: Return a filtered version of the data
- `fillna()`: Return a copy of the data with missing values filled or imputed

We will conclude this section with a brief exploration and demonstration of these routines.

5.1 Detecting null values

Pandas data structures have two useful methods for detecting null data: `isnull()` and `notnull()`. Either one will return a Boolean mask over the data. For example:

```
[22]: df = pd.DataFrame(np.array([[1.5, np.nan, 2.2],
                                   [2.0, 3.5, 5.3],
                                   [np.nan, 4.1, 6.2]]),
                        columns=['col1', 'col2', 'col3'],
                        index=['a', 'b', 'c'])

df
```

```
[22]:   col1  col2  col3
a    1.5   NaN   2.2
b    2.0   3.5   5.3
c    NaN   4.1   6.2
```

```
[23]: df.isnull()
```

```
[23]:   col1  col2  col3
a  False   True  False
b  False  False  False
c   True  False  False
```

The `isnull()` and `notnull()` methods produce similar Boolean results for `DataFrames`.

5.2 Dropping null values

In addition to the masking used before, there are the convenience methods, `dropna()` (which removes NA values) and `fillna()` (which fills in NA values). We cannot drop single values from a `DataFrame`; we can only drop full rows or full columns. Depending on the application, you might want one or the other, so `dropna()` gives a number of options for a `DataFrame`.

By default, `dropna()` will drop all rows in which *any* null value is present:

```
[24]: df.dropna()
```

```
[24]:   col1  col2  col3
b    2.0   3.5   5.3
```

Alternatively, you can drop NA values along a different axis; `axis=1` drops all columns containing a null value:

```
[25]: df.dropna(axis='columns')
```

```
[25]:   col3
a    2.2
b    5.3
c    6.2
```

5.3 Filling null values

Sometimes rather than dropping NA values, you'd rather replace them with a valid value. This value might be a single number like zero, or it might be some sort of imputation or interpolation from the good values. You could do this in-place using the `isnull()` method as a mask, but

because it is such a common operation Pandas provides the `fillna()` method, which returns a copy of the array with the null values replaced.

Consider the following `DataFrame`:

[26]:

```
df
```

```
[26]:      col1  col2  col3
a      1.5   NaN   2.2
b      2.0   3.5   5.3
c      NaN   4.1   6.2
```

For `DataFrames`, we can specify an axis along which the fills take place:

[27]: `df.fillna(method='ffill', axis=1)`

```
[27]:      col1  col2  col3
a      1.5   1.5   2.2
b      2.0   3.5   5.3
c      NaN   4.1   6.2
```

6 Combining Datasets: Concat

Some of the most interesting studies of data come from combining different data sources. These operations can involve anything from very straightforward concatenation of two different datasets, to more complicated database-style joins and merges that correctly handle any overlaps between the datasets. `Series` and `DataFrames` are built with this type of operation in mind, and Pandas includes functions and methods that make this sort of data wrangling fast and straightforward.

Pandas has a function, `pd.concat()`, which has a similar syntax to `np.concatenate` but contains a number of options that we'll discuss momentarily:

Signature in Pandas v0.18

```
pd.concat(objs, axis=0, join='outer', join_axes=None, ignore_index=False,
          keys=None, levels=None, names=None, verify_integrity=False,
          copy=True)
```

`pd.concat()` can be used for a simple concatenation of `Series` or `DataFrame` objects, just as `np.concatenate()` can be used for simple concatenations of arrays:

```
[28]: df1 = pd.DataFrame([[ 'a', 1], [ 'b', 2]],
                        columns=[ 'col1', 'col2'])
      df2 = pd.DataFrame([[ 'c', 3], [ 'd', 4]],
                        columns=[ 'col1', 'col2'])
      display('df1', 'df2', 'pd.concat([df1, df2])')
```

```
[28]: df1
      col1  col2
0      a      1
1      b      2

df2
      col1  col2
```

```
0    c    3
1    d    4
```

```
pd.concat([df1, df2])
   col1  col2
0     a     1
1     b     2
0     c     3
1     d     4
```

By default, the concatenation takes place row-wise within the `DataFrame` (i.e., `axis=0`). Like `np.concatenate`, `pd.concat` allows specification of an axis along which concatenation will take place. Consider the following example:

```
[29]: df3 = pd.DataFrame([[ 'a', 1], [ 'b', 2]],
                        columns=[ 'col1', 'col2'])
      df4 = pd.DataFrame([[ 'c', 3], [ 'd', 4]],
                        columns=[ 'col3', 'col4'])
      display('df3', 'df4', "pd.concat([df3, df4], axis=1)")
```

```
[29]: df3
   col1  col2
0     a     1
1     b     2

df4
   col3  col4
0     c     3
1     d     4

pd.concat([df3, df4], axis=1)
   col1  col2  col3  col4
0     a     1     c     3
1     b     2     d     4
```

We combine `DataFrame` objects horizontally (or vertically) along the x axis by passing in `axis=1` (or `axis=0`)

6.1 Ignoring the index

Sometimes the index itself does not matter, and you would prefer it to simply be ignored. This option can be specified using the `ignore_index` flag. With this set to true, the concatenation will create a new integer index for the resulting `DataFrame`:

```
[30]: display('df1', 'df2', 'pd.concat([df1, df2])', 'pd.concat([df1, df2],
      ↪ignore_index=True)')
```

```
[30]: df1
   col1  col2
0     a     1
1     b     2
```

```
df2
  col1  col2
0    c     3
1    d     4

pd.concat([df1, df2])
  col1  col2
0    a     1
1    b     2
0    c     3
1    d     4

pd.concat([df1, df2], ignore_index=True)
  col1  col2
0    a     1
1    b     2
2    c     3
3    d     4
```

6.2 Concatenation with joins

In the simple examples we just looked at, we were mainly concatenating `DataFrames` with shared column names. In practice, data from different sources might have different sets of column names, and `pd.concat` offers several options in this case. Consider the concatenation of the following two `DataFrames`, which have some (but not all!) columns in common:

```
[31]: df5 = pd.DataFrame(['a', 1], ['b', 2],
                        columns=['col1', 'col2'])
df6 = pd.DataFrame(['c', 3], ['d', 4],
                  columns=['col2', 'col3'])
display('df5', 'df6', 'pd.concat([df5, df6])')
```

```
[31]: df5
  col1  col2
0    a     1
1    b     2

df6
  col2  col3
0    c     3
1    d     4

pd.concat([df5, df6])
  col1  col2  col3
0    a     1   NaN
1    b     2   NaN
0  NaN    c   3.0
1  NaN    d   4.0
```

By default, the entries for which no data is available are filled with NA values. To change this, we can specify one of several options for the `join` and `join_axes` parameters of the `concatenate` function. By default, the join is a union of the input columns (`join='outer'`), but we can

change this to an intersection of the columns using `join='inner'`:

```
[32]: display('df5', 'df6',
            "pd.concat([df5, df6], join='inner')")
```

```
[32]: df5
      col1  col2
0       a     1
1       b     2

df6
      col2  col3
0       c     3
1       d     4

pd.concat([df5, df6], join='inner')
      col2
0       1
1       2
0       c
1       d
```

The combination of options of the `pd.concat` function allows a wide range of possible behaviors when joining two datasets; keep these in mind as you use these tools for your own data.

7 Combining Datasets: Merge

One essential feature offered by Pandas is its high-performance, in-memory join and merge operations. If you have ever worked with databases, you should be familiar with this type of data interaction. The main interface for this is the `pd.merge` function, and we'll see few examples of how this can work in practice. The main difference between `pd.merge()` and `pd.concat()` is:

- `pd.merge()` is used to combine two (or more) dataframes on the basis of values of common columns.
- `pd.concat()` is used to append one (or more) dataframes one below the other (or sideways)

7.1 Categories of Joins

The `pd.merge()` function implements a number of types of joins: the *one-to-one*, *many-to-one*, and *many-to-many* joins. All three types of joins are accessed via an identical call to the `pd.merge()` interface; the type of join performed depends on the form of the input data. Here we will show simple examples of the three types of merges, and discuss detailed options further below.

7.1.1 One-to-one joins

Perhaps the simplest type of merge expression is the one-to-one join, which is in many ways very similar to the column-wise concatenation. As a concrete example, consider the following two `DataFrames` which contain information on several employees in a company:

```
[33]: df1 = pd.DataFrame({'employee': ['Bob', 'Jake', 'Lisa', 'Sue'],
                        'group': ['Accounting', 'Engineering', 'Engineering', 'HR']})
```

```
df2 = pd.DataFrame({'employee': ['Lisa', 'Bob', 'Jake', 'Sue'],
                    'hire_date': [2004, 2008, 2012, 2014]})
display('df1', 'df2')
```

```
[33]: df1
      employee      group
0        Bob  Accounting
1        Jake  Engineering
2        Lisa  Engineering
3         Sue           HR

df2
      employee  hire_date
0        Lisa       2004
1         Bob       2008
2        Jake       2012
3         Sue       2014
```

To combine this information into a single `DataFrame`, we can use the `pd.merge()` function:

```
[34]: df3 = pd.merge(df1, df2)
df3
```

```
[34]:   employee      group  hire_date
0        Bob  Accounting       2008
1        Jake  Engineering       2012
2        Lisa  Engineering       2004
3         Sue           HR       2014
```

The `pd.merge()` function recognizes that each `DataFrame` has an “employee” column, and automatically joins using this column as a key. The result of the merge is a new `DataFrame` that combines the information from the two inputs. Notice that the order of entries in each column is not necessarily maintained: in this case, the order of the “employee” column differs between `df1` and `df2`, and the `pd.merge()` function correctly accounts for this. Additionally, keep in mind that the merge in general discards the index, except in the special case of merges by index (see the `left_index` and `right_index` keywords, discussed momentarily).

7.1.2 Many-to-one joins

Many-to-one joins are joins in which one of the two key columns contains duplicate entries. For the many-to-one case, the resulting `DataFrame` will preserve those duplicate entries as appropriate. Consider the following example of a many-to-one join:

```
[35]: df4 = pd.DataFrame({'group': ['Accounting', 'Engineering', 'HR'],
                          'supervisor': ['Carly', 'Guido', 'Steve']})
display('df3', 'df4', 'pd.merge(df3, df4)')
```

```
[35]: df3
      employee      group  hire_date
0        Bob  Accounting       2008
1        Jake  Engineering       2012
2        Lisa  Engineering       2004
```

```
3      Sue      HR      2014
```

```
df4
```

```
      group supervisor
0  Accounting      Carly
1  Engineering      Guido
2         HR      Steve
```

```
pd.merge(df3, df4)
```

```
  employee      group  hire_date supervisor
0      Bob  Accounting      2008      Carly
1      Jake  Engineering      2012      Guido
2      Lisa  Engineering      2004      Guido
3      Sue      HR      2014      Steve
```

The resulting `DataFrame` has an additional column with the “supervisor” information, where the information is repeated in one or more locations as required by the inputs.

7.1.3 Many-to-many joins

Many-to-many joins are a bit confusing conceptually, but are nevertheless well defined. If the key column in both the left and right array contains duplicates, then the result is a many-to-many merge. This will be perhaps most clear with a concrete example. Consider the following, where we have a `DataFrame` showing one or more skills associated with a particular group. By performing a many-to-many join, we can recover the skills associated with any individual person:

```
[36]: df5 = pd.DataFrame({'group': ['Accounting', 'Accounting',
                                   'Engineering', 'Engineering', 'HR', 'HR'],
                          'skills': ['math', 'spreadsheets', 'coding', 'linux',
                                    'spreadsheets', 'organization']})
display('df1', 'df5', "pd.merge(df1, df5)")
```

```
[36]: df1
```

```
  employee      group
0      Bob  Accounting
1      Jake  Engineering
2      Lisa  Engineering
3      Sue      HR
```

```
df5
```

```
      group      skills
0  Accounting      math
1  Accounting  spreadsheets
2  Engineering      coding
3  Engineering      linux
4         HR  spreadsheets
5         HR  organization
```

```
pd.merge(df1, df5)
```

```
  employee      group      skills
0      Bob  Accounting      math
1      Bob  Accounting  spreadsheets
```

2	Jake	Engineering	coding
3	Jake	Engineering	linux
4	Lisa	Engineering	coding
5	Lisa	Engineering	linux
6	Sue	HR	spreadsheets
7	Sue	HR	organization

These three types of joins can be used with other Pandas tools to implement a wide array of functionality. But in practice, datasets are rarely as clean as the one we're working with here. In the following section we'll consider some of the options provided by `pd.merge()` that enable you to tune how the join operations work.

7.2 Specification of the Merge Key

We've already seen the default behavior of `pd.merge()`: it looks for one or more matching column names between the two inputs, and uses this as the key. However, often the column names will not match so nicely, and `pd.merge()` provides a variety of options for handling this.

7.2.1 The `on` keyword

Most simply, you can explicitly specify the name of the key column using the `on` keyword, which takes a column name or a list of column names:

```
[37]: display('df1', 'df2', "pd.merge(df1, df2, on='employee')")
```

```
[37]: df1
   employee  group
0      Bob  Accounting
1      Jake  Engineering
2      Lisa  Engineering
3      Sue      HR

df2
   employee  hire_date
0      Lisa      2004
1       Bob      2008
2      Jake      2012
3      Sue      2014

pd.merge(df1, df2, on='employee')
   employee  group  hire_date
0      Bob  Accounting      2008
1      Jake  Engineering      2012
2      Lisa  Engineering      2004
3      Sue      HR      2014
```

This option works only if both the left and right DataFrames have the specified column name.

7.2.2 The `left_on` and `right_on` keywords

At times you may wish to merge two datasets with different column names; for example, we may have a dataset in which the employee name is labeled as “name” rather than “employee”. In this case, we can use the `left_on` and `right_on` keywords to specify the two column names:

```
[38]: df3 = pd.DataFrame({'name': ['Bob', 'Jake', 'Lisa', 'Sue'],
                        'salary': [70000, 80000, 120000, 90000]})
display('df1', 'df3', 'pd.merge(df1, df3, left_on="employee",
    →right_on="name")')
```

```
[38]: df1
      employee      group
0         Bob  Accounting
1         Jake  Engineering
2         Lisa  Engineering
3          Sue           HR

df3
   name  salary
0   Bob   70000
1  Jake   80000
2  Lisa  120000
3   Sue   90000

pd.merge(df1, df3, left_on="employee", right_on="name")
   employee      group  name  salary
0         Bob  Accounting  Bob   70000
1         Jake  Engineering  Jake   80000
2         Lisa  Engineering  Lisa  120000
3          Sue           HR   Sue   90000
```

The result has a redundant column that we can drop if desired—for example, by using the `drop()` method of DataFrames:

```
[39]: pd.merge(df1, df3, left_on="employee", right_on="name").drop('name', axis=1)
```

```
[39]:   employee      group  salary
0         Bob  Accounting   70000
1         Jake  Engineering   80000
2         Lisa  Engineering  120000
3          Sue           HR   90000
```

All of these options also work with multiple indices and/or multiple columns; the interface for this behavior is very intuitive. For more information on this, see the [“Merge, Join, and Concatenate”](#) section of the Pandas documentation.

8 Aggregation and Grouping

An essential piece of analysis of large data is efficient summarization: computing aggregations like `sum()`, `mean()`, `median()`, `min()`, and `max()`, in which a single number gives insight into the nature of a potentially large dataset. In this section, we’ll explore aggregations in Pandas, from simple operations akin to what we’ve seen on NumPy arrays, to more sophisticated operations based on the concept of a `groupby`.

8.1 Simple Aggregation in Pandas

Here we will use the Planets dataset, available via the [Seaborn package](#). It gives information on planets that astronomers have discovered around other stars (known as *extrasolar planets* or *exoplanets* for short). It can be downloaded with a simple Seaborn command:

```
[40]: import seaborn as sns
planets = sns.load_dataset('planets')
planets.shape
```

```
[40]: (1035, 6)
```

```
[41]: planets.head()
```

```
[41]:
```

	method	number	orbital_period	mass	distance	year
0	Radial Velocity	1	269.300	7.10	77.40	2006
1	Radial Velocity	1	874.774	2.21	56.95	2008
2	Radial Velocity	1	763.000	2.60	19.84	2011
3	Radial Velocity	1	326.030	19.40	110.62	2007
4	Radial Velocity	1	516.220	10.50	119.47	2009

This has some details on the 1,000+ extrasolar planets discovered up to 2014.

Pandas `DataFrame` includes all of the common aggregates; in addition, there is a convenience method `describe()` that computes several common aggregates for each column and returns the result. Let's use this on the Planets data, for now dropping rows with missing values:

```
[42]: planets.dropna().describe()
```

```
[42]:
```

	number	orbital_period	mass	distance	year
count	498.000000	498.000000	498.000000	498.000000	498.000000
mean	1.73494	835.778671	2.509320	52.068213	2007.377510
std	1.17572	1469.128259	3.636274	46.596041	4.167284
min	1.000000	1.328300	0.003600	1.350000	1989.000000
25%	1.000000	38.272250	0.212500	24.497500	2005.000000
50%	1.000000	357.000000	1.245000	39.940000	2009.000000
75%	2.000000	999.600000	2.867500	59.332500	2011.000000
max	6.000000	17337.500000	25.000000	354.000000	2014.000000

This can be a useful way to begin understanding the overall properties of a dataset. For example, we see in the `year` column that although exoplanets were discovered as far back as 1989, half of all known exoplanets were not discovered until 2010 or after. This is largely thanks to the *Kepler* mission, which is a space-based telescope specifically designed for finding eclipsing planets around other stars.

The following table summarizes some other built-in Pandas aggregations:

Aggregation	Description
<code>count()</code>	Total number of items
<code>first()</code> , <code>last()</code>	First and last item
<code>mean()</code> , <code>median()</code>	Mean and median
<code>min()</code> , <code>max()</code>	Minimum and maximum
<code>std()</code> , <code>var()</code>	Standard deviation and variance
<code>mad()</code>	Mean absolute deviation

Aggregation	Description
<code>prod()</code>	Product of all items
<code>sum()</code>	Sum of all items

These are all methods of `DataFrame` and `Series` objects.

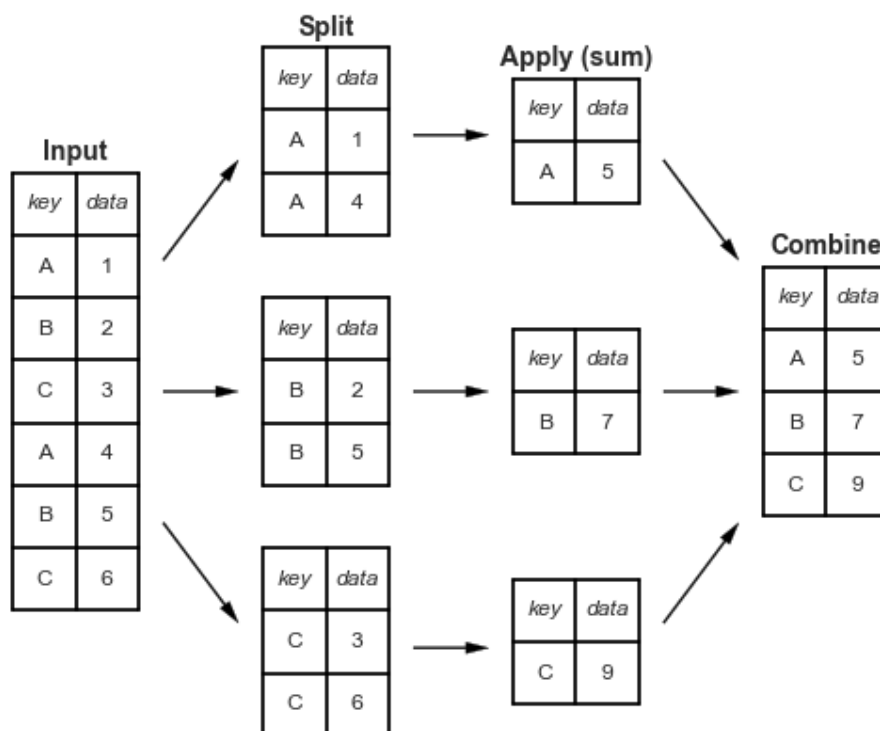
To go deeper into the data, however, simple aggregates are often not enough. The next level of data summarization is the `groupby` operation, which allows you to quickly and efficiently compute aggregates on subsets of data.

8.2 GroupBy: Split, Apply, Combine

Simple aggregations can give you a flavor of your dataset, but often we would prefer to aggregate conditionally on some label or index: this is implemented in the so-called `groupby` operation. The name “group by” comes from a command in the SQL database language, but it is perhaps more illuminative to think of it in the terms first coined by Hadley Wickham of Rstats fame: *split, apply, combine*.

8.2.1 Split, apply, combine

A canonical example of this split-apply-combine operation, where the “apply” is a summation aggregation, is illustrated in this figure:



This makes clear what the `groupby` accomplishes:

- The *split* step involves breaking up and grouping a `DataFrame` depending on the value of the specified key.
- The *apply* step involves computing some function, usually an aggregate, transformation, or filtering, within the individual groups.

- The *combine* step merges the results of these operations into an output array.

While this could certainly be done manually using some combination of the masking, aggregation, and merging commands covered earlier, an important realization is that *the intermediate splits do not need to be explicitly instantiated*. Rather, the `GroupBy` can (often) do this in a single pass over the data, updating the sum, mean, count, min, or other aggregate for each group along the way. The power of the `GroupBy` is that it abstracts away these steps: the user need not think about *how* the computation is done under the hood, but rather thinks about the *operation as a whole*.

As a concrete example, let's take a look at using Pandas for the computation shown in this diagram. We'll start by creating the input `DataFrame`:

```
[43]: df = pd.DataFrame({'key': ['A', 'B', 'C', 'A', 'B', 'C'],
                        'data1': range(6),
                        'data2': range(5,11)}, columns=['key', 'data1', 'data2'])
df
```

```
[43]:   key  data1  data2
0    A      0      5
1    B      1      6
2    C      2      7
3    A      3      8
4    B      4      9
5    C      5     10
```

The most basic split-apply-combine operation can be computed with the `groupby()` method of `DataFrames`, passing the name of the desired key column:

```
[44]: df.groupby('key')
```

```
[44]: <pandas.core.groupby.generic.DataFrameGroupBy object at 0x00000175A9DCB730>
```

Notice that what is returned is not a set of `DataFrames`, but a `DataFrameGroupBy` object. This object is where the magic is: you can think of it as a special view of the `DataFrame`, which is poised to dig into the groups but does no actual computation until the aggregation is applied. This “lazy evaluation” approach means that common aggregates can be implemented very efficiently in a way that is almost transparent to the user.

To produce a result, we can apply an aggregate to this `DataFrameGroupBy` object, which will perform the appropriate apply/combine steps to produce the desired result:

```
[45]: df.groupby('key').sum()
```

```
[45]:   data1  data2
key
A      3     13
B      5     15
C      7     17
```

```
[46]: df.groupby('key')['data1'].sum()
```

```
[46]: key
A      3
```

```
B    5
C    7
Name: data1, dtype: int64
```

The `sum()` method is just one possibility here; you can apply virtually any common Pandas or NumPy aggregation function, as well as virtually any valid `DataFrame` operation, as we will see in the following discussion.

9 Reading and Writing Files

We can load a csv file into a `DataFrame` object (using `read_csv`). Here we will consider an example of some data about US states and their populations. The data are available in the files [state_population.csv](#) and [state_abbrevs.csv](#).

```
[47]: pop = pd.read_csv('data/state_population.csv')
      abbrevs = pd.read_csv('data/state_abbrevs.csv')
      display('pop.head()', 'abbrevs.head()')
```

```
[47]: pop.head()
      state/region  ages  year  population
0             AL  under18  2012    1117489.0
1             AL    total  2012    4817528.0
2             AL  under18  2010    1130966.0
3             AL    total  2010    4785570.0
4             AL  under18  2011    1125763.0
```

```
abbrevs.head()
      state abbreviation
0     Alabama          AL
1     Alaska           AK
2     Arizona          AZ
3  Arkansas           AR
4  California          CA
```

```
[48]: merged = pd.merge(pop, abbrevs, left_on='state/region',
      →right_on='abbreviation')
      merged = merged.drop(columns=['abbreviation']) # drop duplicate info
      merged.head()
```

```
[48]: state/region  ages  year  population  state
0             AL  under18  2012    1117489.0  Alabama
1             AL    total  2012    4817528.0  Alabama
2             AL  under18  2010    1130966.0  Alabama
3             AL    total  2010    4785570.0  Alabama
4             AL  under18  2011    1125763.0  Alabama
```

We can store a `DataFrame` object into a csv file (using `to_csv`)

```
[49]: merged.to_csv('my_dataframe.csv')
```